


Automated verification of systems software with Serval


Luke Nelson, Emina Torlak, Xi Wang
Paul G. Allen School of Computer Science & Engineering
University of Washington




Outline

- Part 1: Overview
- Part 2: Serval
- Part 3: Rosette

Part I: overview



specification



proof



implementation

Overview outline

- History
- Example: specification & verification
- Project organization

Pioneer efforts of OS kernel verification

- Examples of earlier efforts
 - UCLA Unix security kernel
 - Kit
- seL4 microkernel
 - first functional correctness proof (2009)
 - 10,000 lines of C and assembly

Examples in recent SOSPP/OSDI

AtomFS

Notary

Perennial

Serval

Vigor

DFSCQ

CSPEC

Ironclad

FSCQ

CertiKOS

Hyperkernel

Nickel

Jitk

IronFleet

Yggdrasil

Komodo

SFSCQ



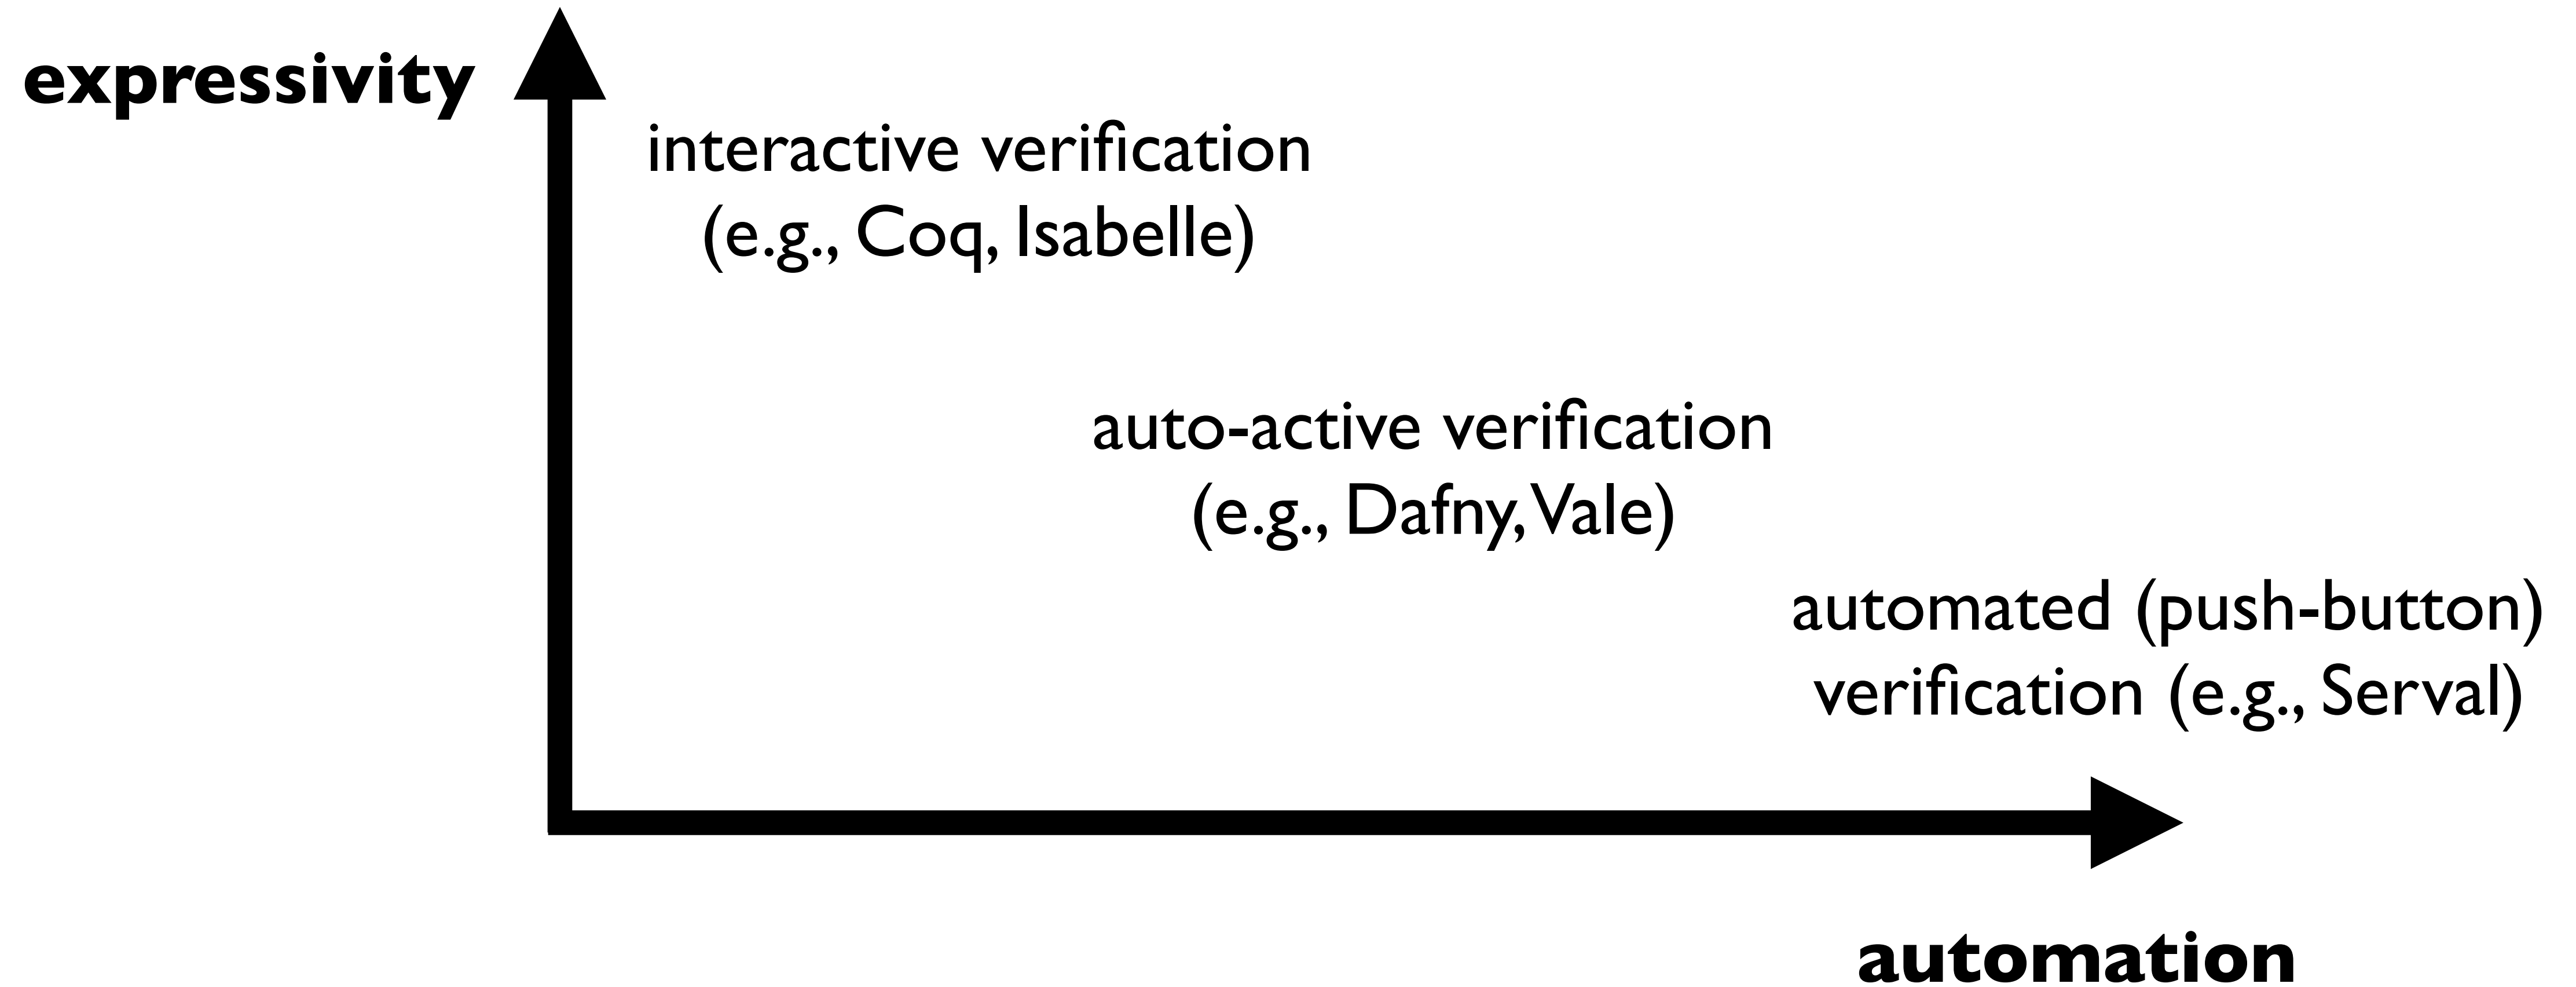
Types of systems verified

- Hardware
- OS kernels & security monitors
- File systems
- Distributed systems
- Networking
- Applications

Properties & specifications

- State-machine refinement
- Noninterference
- Crash safety
- Determinism
- Linearizability
- Liveness

Verification methodologies & tools



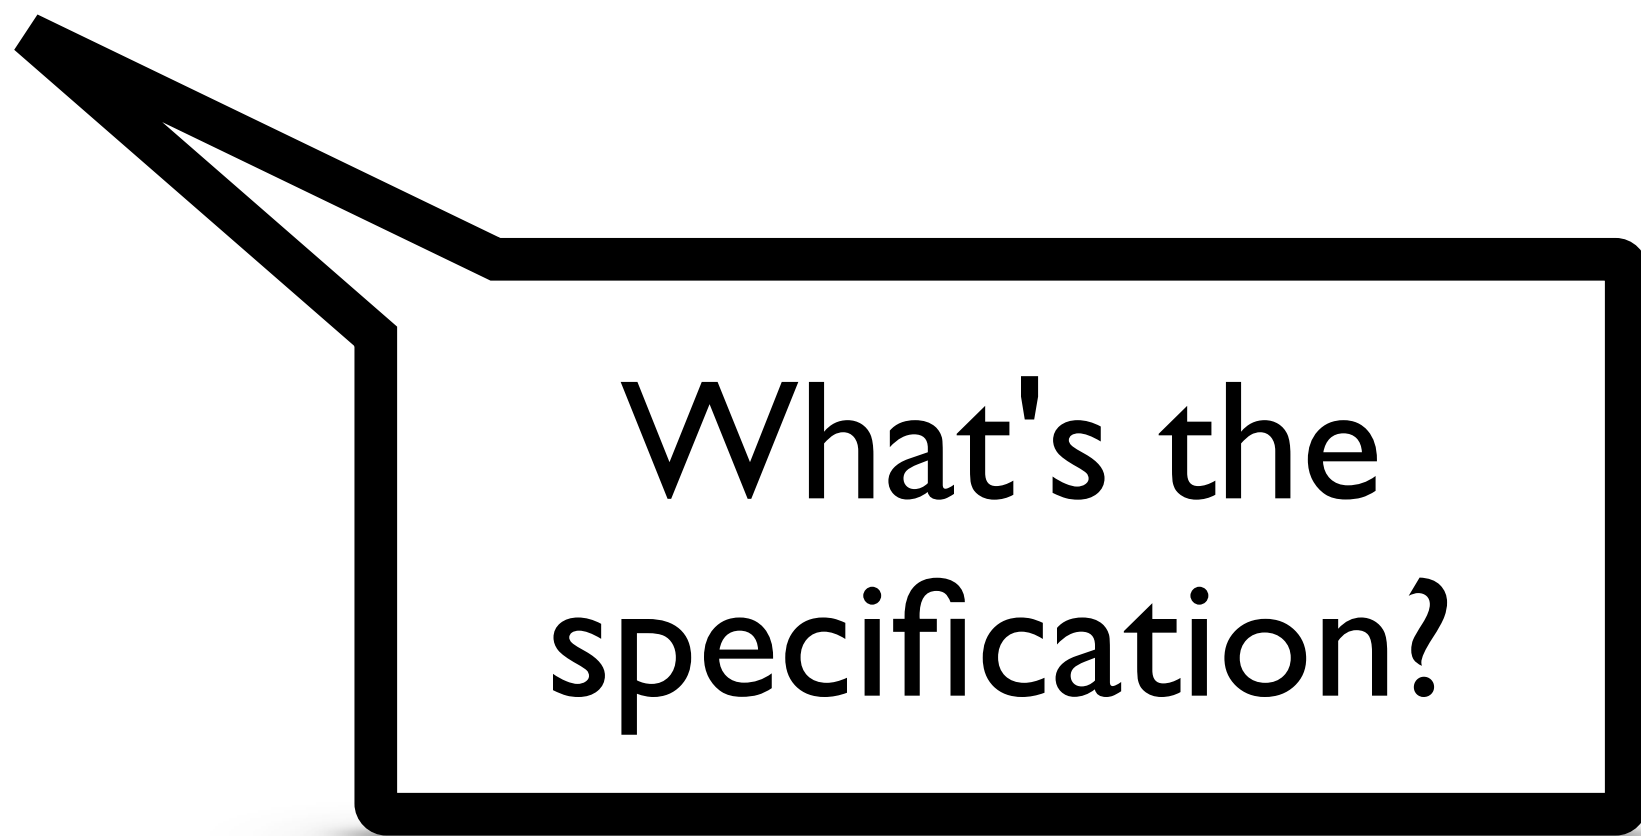
Overview outline

- History
- **Example: specification & verification**
- Project organization

Example: verify a 16-bit integer encoder/decoder

```
fn encode(x: u16) -> (u8, u8) {  
    (x as u8, (x >> 8) as u8)  
}
```

```
fn decode(b: (u8, u8)) -> u16 {  
    (b.0 as u16) | ((b.1 as u16) << 8)  
}
```



What's the
specification?

Specification example

- Theorem: forall 16-bit x, $\text{decode}(\text{encode}(x)) == x$
- Any encoded 16-bit integer can be decoded back
- What bugs cannot be captured by the specification?

Specification example

- Theorem: forall 16-bit x, $\text{decode}(\text{encode}(x)) == x$
- Any encoded 16-bit integer can be decoded back
- What bugs cannot be captured by the specification?
 - Correctness of encode (or decode) alone
 - Whether any two bytes can be decoded and encoded back (e.g., parser)

Specification alternatives

- Theorem: forall 16-bit x , $\text{decode}(\text{encode}(x)) == x$
- Any encoded 16-bit integer can be decoded back
- Alternative spec: model little endianness for decode/encode
- Alternative spec: memory safety for encode/decode

Summary of specification

- Specification is key to verification
 - Theorem & model of environment
 - Any bugs in specification can invalidate guarantees
- Trade-offs
 - Expressiveness: does the spec prevent the intended bugs?
 - Simplicity: is the spec easy to audit?

Verification

forall 16-bit x, decode(encode(x)) == x

```
fn encode(x: u16) -> (u8, u8) {  
    (x as u8, (x >> 8) as u8)  
}  
  
fn decode(b: (u8, u8)) -> u16 {  
    (b.0 as u16) | ((b.1 as u16) << 8)  
}
```



specification

implementation

Verify impl against spec 1/2: enumerating

- Theorem: forall 16-bit x , $\text{decode}(\text{encode}(x)) == x$
- Exhaustively enumerate every possible value of x : 0 to 65535
- Easy to automate; hard to scale (if the input space is large)

Verify impl against spec 2/2: rewriting

- Theorem: forall 16-bit x, decode(encode(x)) == x
- Repeatedly apply rewrite rules until true
- Demo: run rustc/llvm

```
fn spec_encode_decode(x: u16) -> bool {  
    decode(encode(x)) == x  
}
```

- Hard to automate; easy to generalize to larger integer types

Summary of verification

- Two basic proof strategies
 - Search in the space of input data
 - Search in the space of rewriting rules
- Usually use a hybrid approach

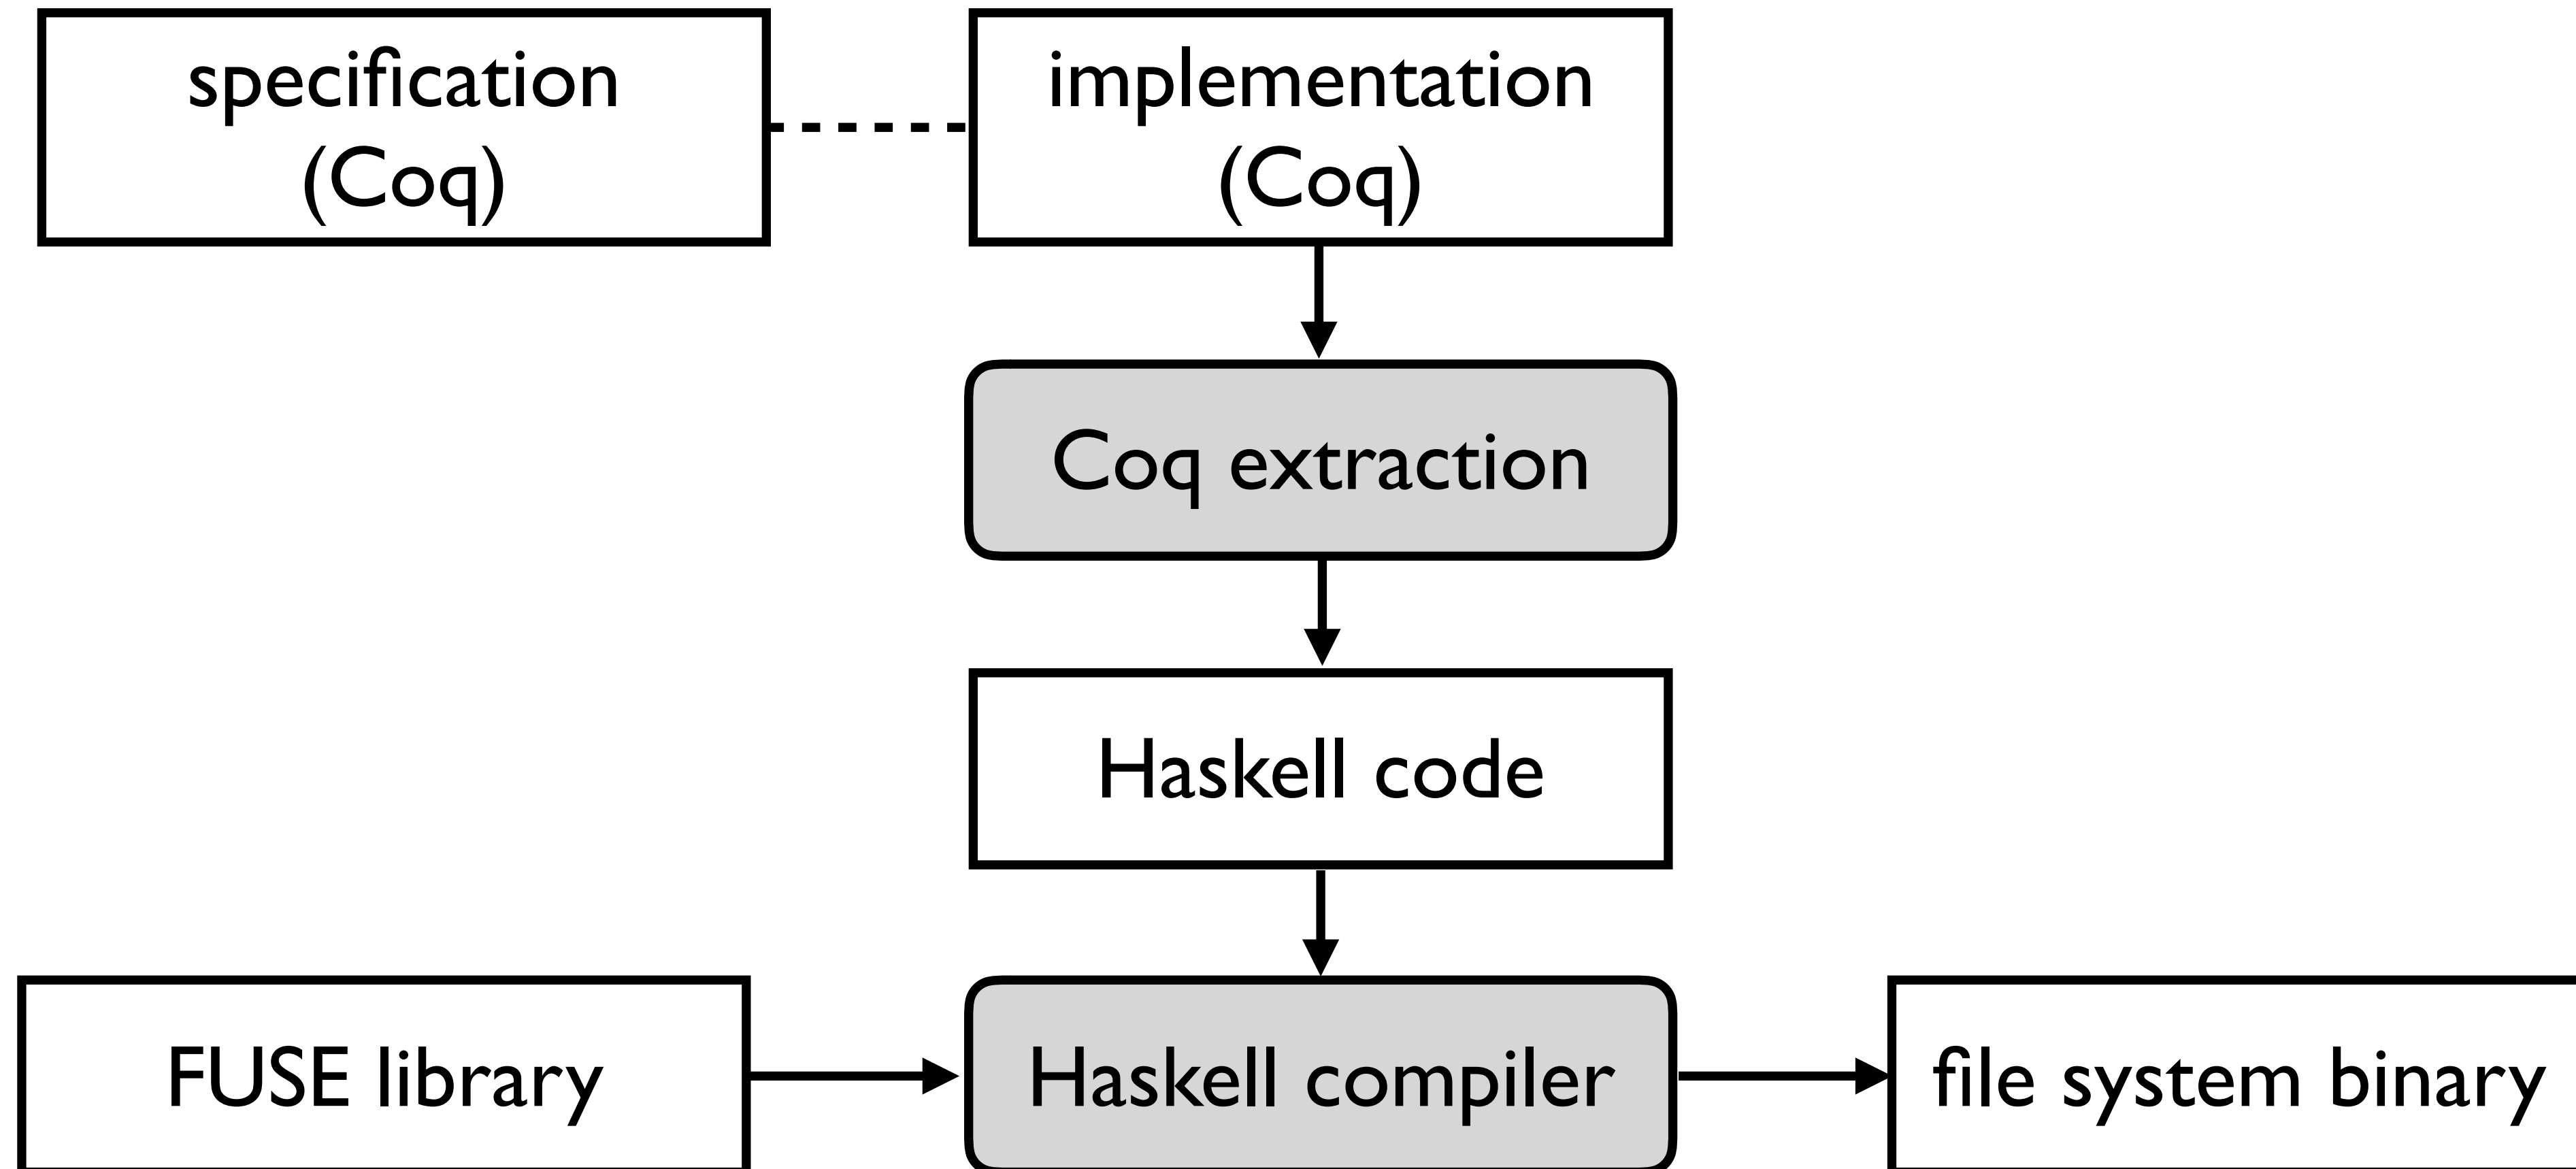
Overview outline

- History
- Example: specification & verification
- **Project organization**

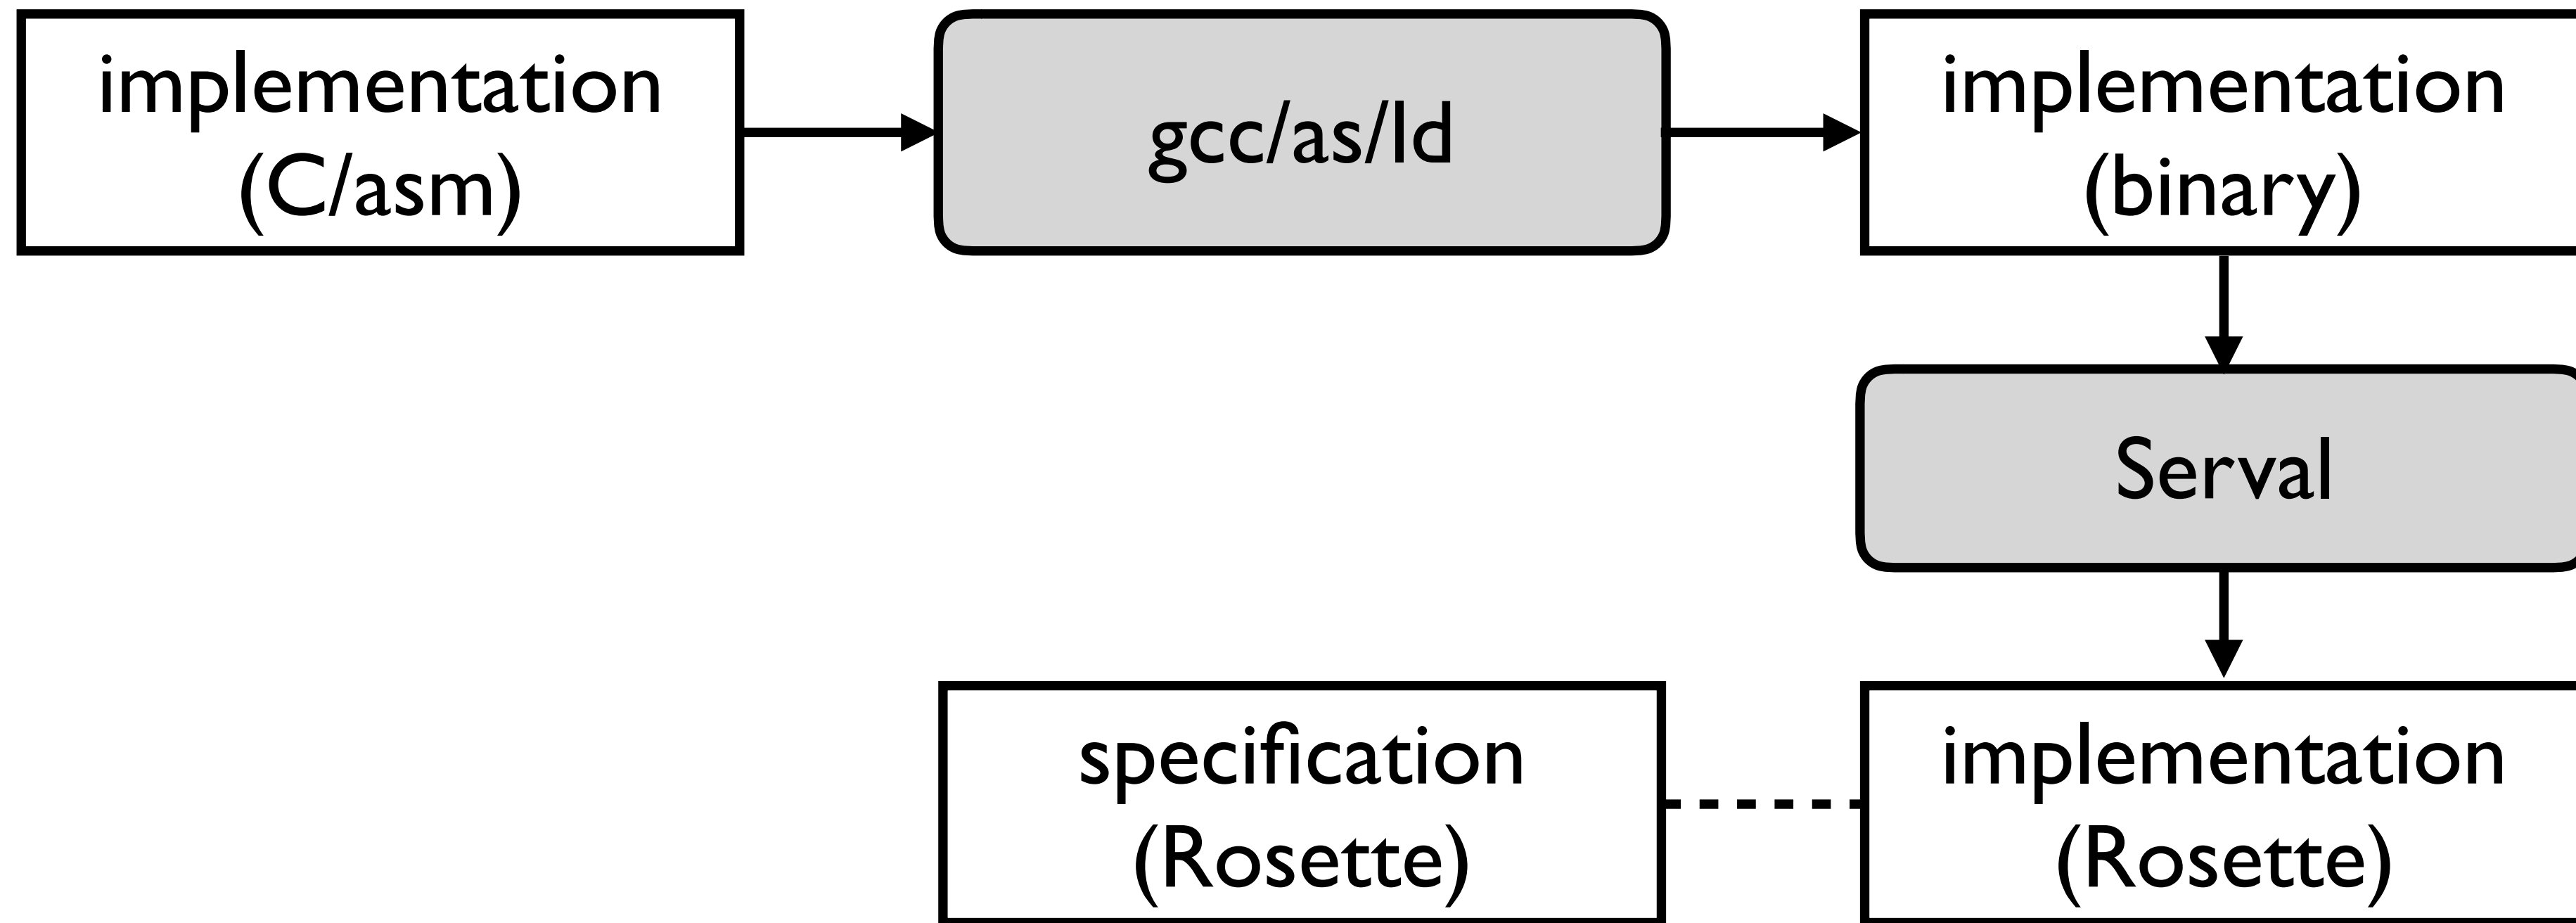
How to organize a verification project

- Extraction
- Parsing
- Verified compilation

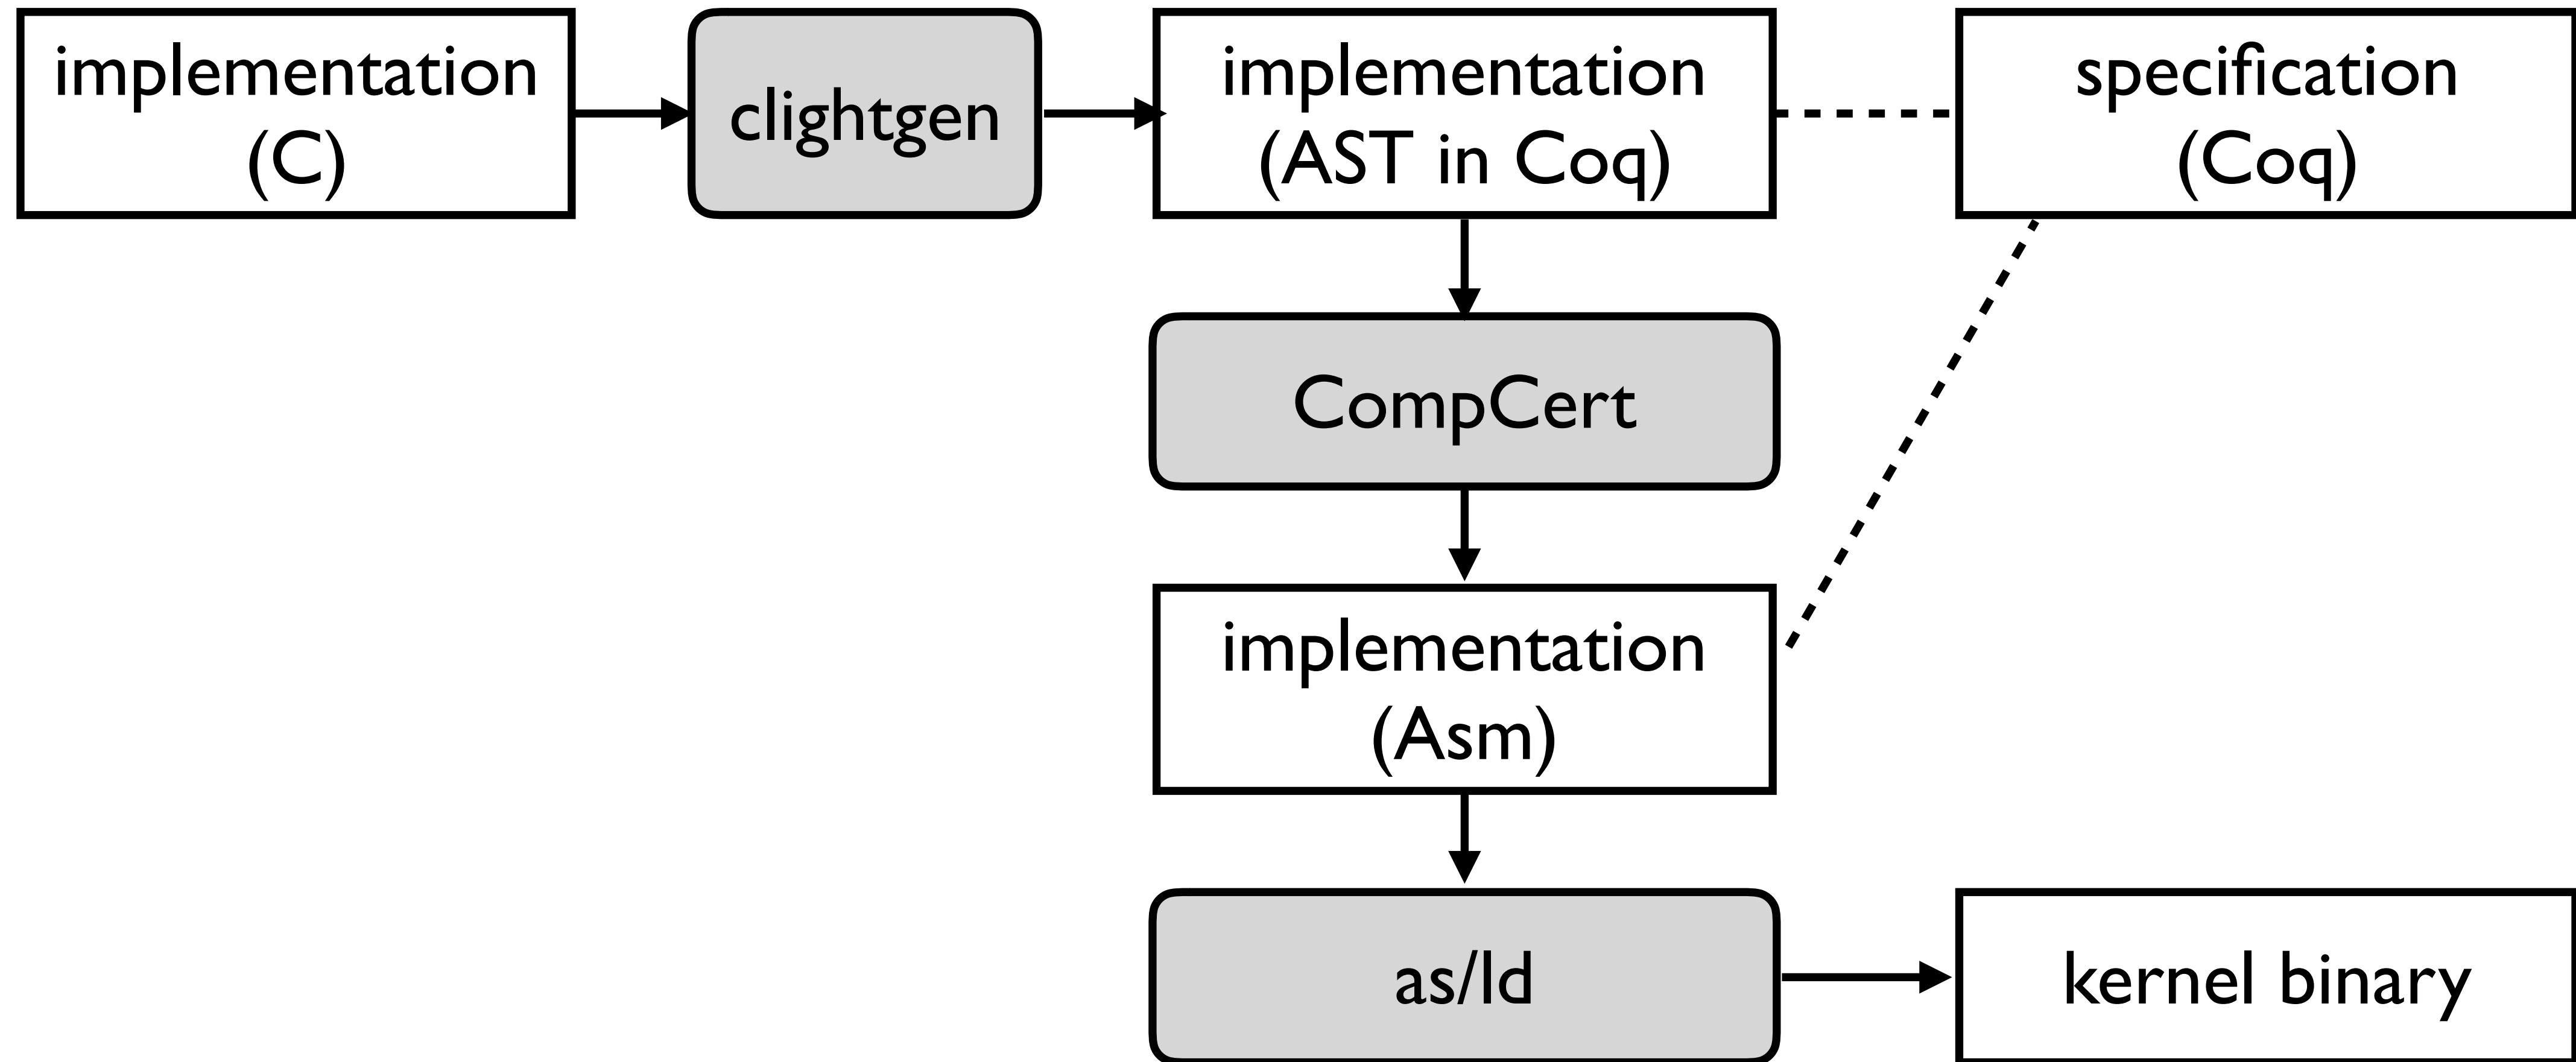
Extraction example: FSCQ



Parsing example: Serval



Verified compilation example: CertiKOS



Summary of Part I

- Verification is effective at eliminating bugs
- Choose the "right" trade-offs for your system
 - Specification & design
 - Verification methodologies
 - Turning implementation to executable code

Part 2: Serval

- Overview of Serval
- A toy monitor
- Example: free of low-level bugs
- Example: functional correctness
- Example: safety properties

Verification stack

System
specification

RISC-V
instructions

x86-32
instructions

LLVM
instructions

BPF
instructions

RISC-V
verifier

x86-32
verifier

LLVM
verifier

BPF verifier



Serval:

Specification library, symbolic optimizations, machine code support



Rosette:

Symbolic evaluation, symbolic profiling, symbolic reflection



SMT solver:

Constraint solving, counterexample generation

Verification stack

System
specification

RISC-V
instructions

x86-32
instructions

LLVM
instructions

BPF
instructions

RISC-V
verifier

x86-32
verifier

LLVM
verifier

BPF verifier



Serval:

Specification library, symbolic optimizations, machine code support



Rosette:

Symbolic evaluation, symbolic profiling, symbolic reflection



SMT solver:

Constraint solving, counterexample generation

Verification stack

System
specification

RISC-V
instructions

x86-32
instructions

LLVM
instructions

BPF
instructions

RISC-V
verifier

x86-32
verifier

LLVM
verifier

BPF verifier

Serval:

Specification library, symbolic optimizations, machine code support

Rosette:

Symbolic evaluation, symbolic profiling, symbolic reflection

SMT solver:

Constraint solving, counterexample generation



Z3

Verification stack

System
specification

RISC-V
instructions

x86-32
instructions

LLVM
instructions

BPF
instructions

RISC-V
verifier

x86-32
verifier

LLVM
verifier

BPF verifier

Serval:

Specification library, symbolic optimizations, machine code support

Rosette:

Symbolic evaluation, symbolic profiling, symbolic reflection

SMT solver:

Constraint solving, counterexample generation



Z3

Verification stack

System
specification

RISC-V
instructions

x86-32
instructions

LLVM
instructions

BPF
instructions

RISC-V
verifier

x86-32
verifier

LLVM
verifier

BPF verifier

Serval:

Specification library, symbolic optimizations, machine code support

Rosette:

Symbolic evaluation, symbolic profiling, symbolic reflection

SMT solver:

Constraint solving, counterexample generation



Z3

Verification stack

System
specification

RISC-V
instructions

x86-32
instructions

LLVM
instructions

BPF
instructions

RISC-V
verifier

x86-32
verifier

LLVM
verifier

BPF verifier



Serval:

Specification library, symbolic optimizations, machine code support



Rosette:

Symbolic evaluation, symbolic profiling, symbolic reflection

Z3

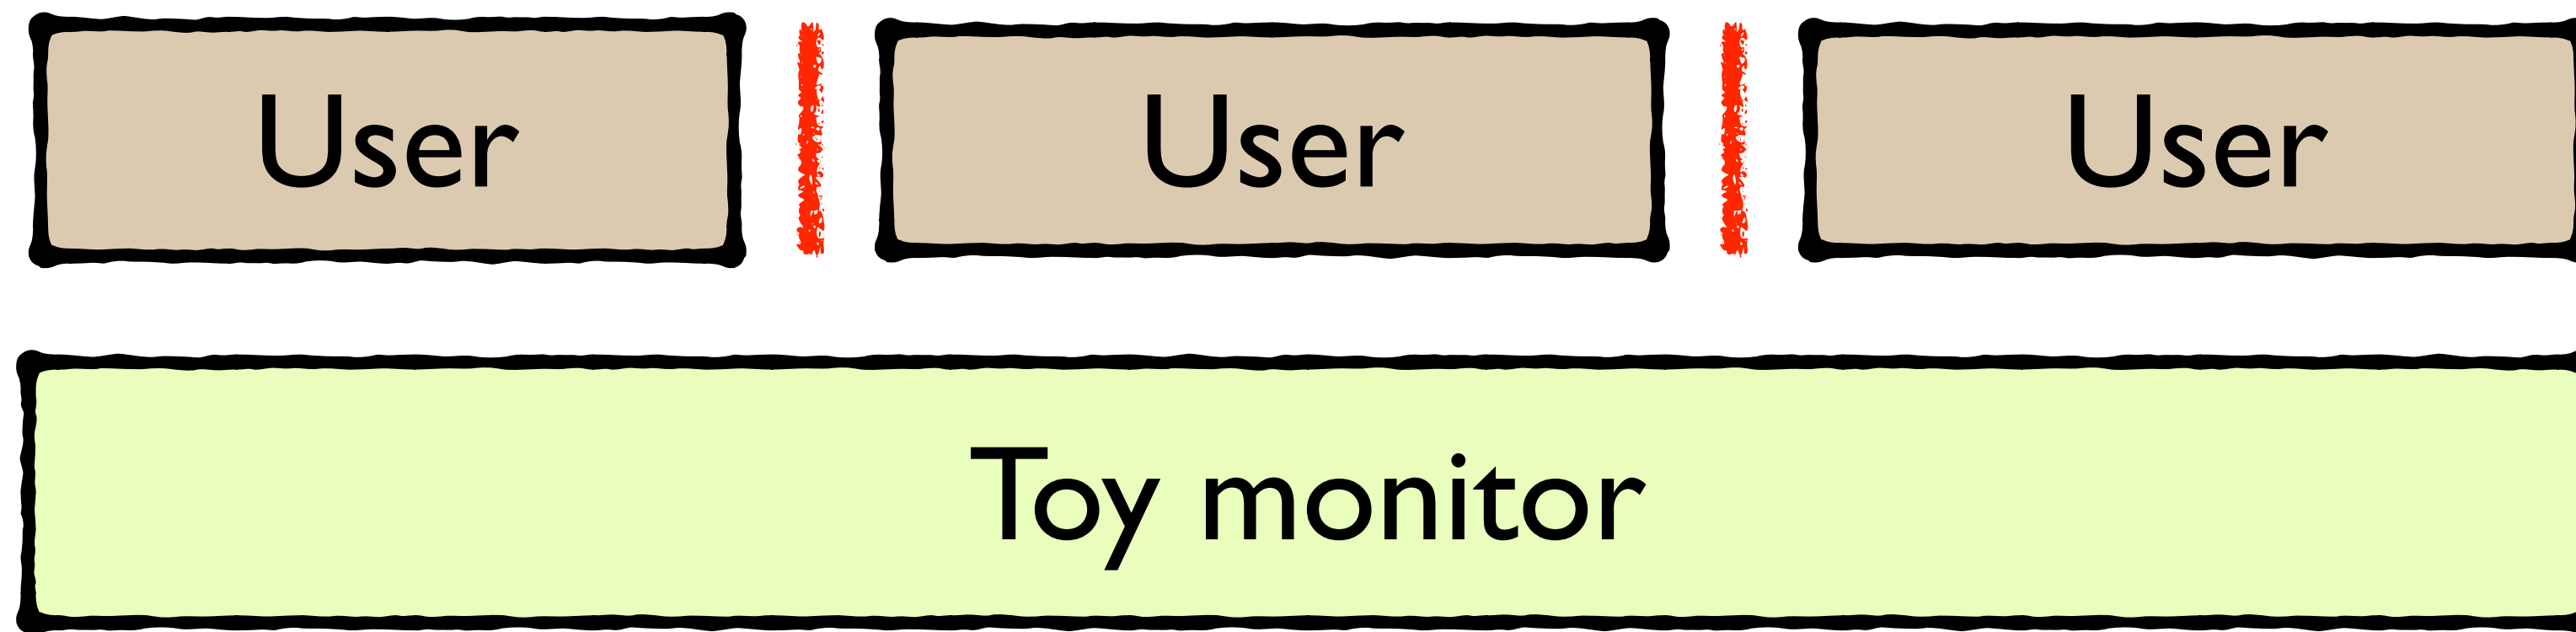
SMT solver:

Constraint solving, counterexample generation


What you can do with Serval

- **Verify functional correctness and safety properties**
- Verify system designs: Keystone
- Verify existing code: BPF JIT in the Linux kernel

A toy monitor




Verifying state-machine refinement



Refinement: both spec and impl state machines move in lock-step

Verifying noninterference (step consistency)



Confidentiality: stepping from two equivalent states results in equivalent states (e.g., parts of the state that cannot be observed by user P cannot affect its execution)